Problem Statement: - Social Network Ad

How to develop an AI solution to personalize advertisements for users based on historical data, and predict whether a user is likely to purchase the advertised product. The AI system should display advertisements only to users with a high probability of making a purchase, in order to target the right customers effectively.

Class 0 – User purchased

Class 1 – User did not purchase

Support – Actual No. of. Samples in Each class

Random Forest Classifier

Code link: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/1.RandomForest_classification.ipynb

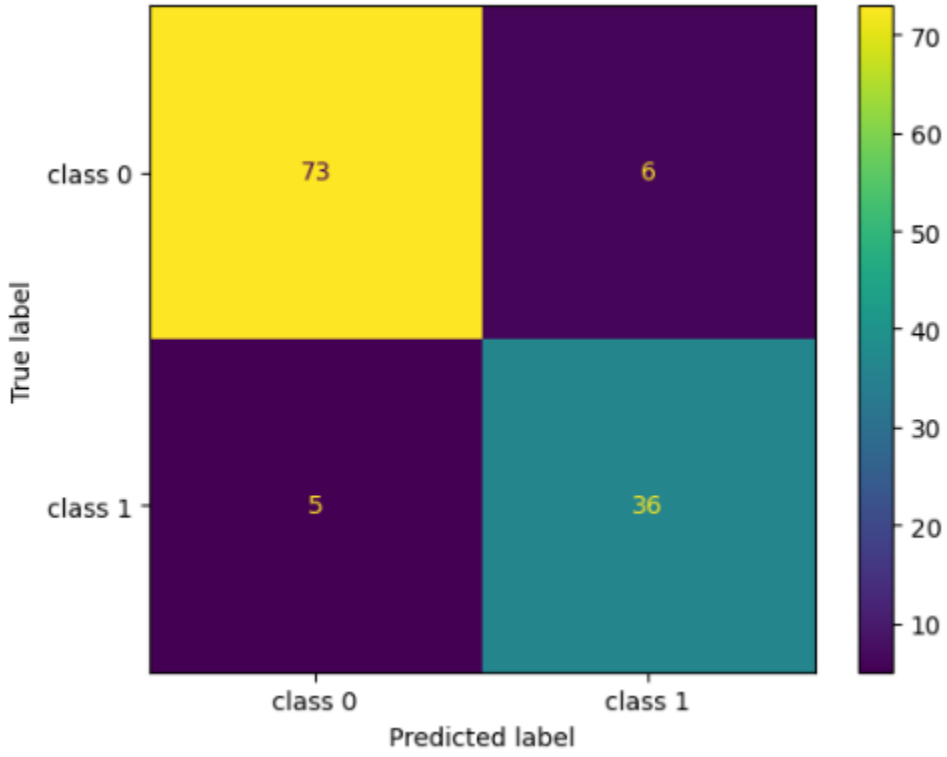

	precision	recall	f1-score	support
0	0.94	0.92	0.93	79
1	0.86	0.88	0.87	41
accuracy			0.91	120
macro avg	0.90	0.90	0.90	120
weighted avg	0.91	0.91	0.91	120

Classification report for random forest classifier

1) What is the overall performance of the model?
   Accuracy : 0.91
2) What is the percentage of correctly classified class 0?
   Recall of class 0 : 0.92
3) What is the percentage of correctly classified class 1?
   Recall of class 1 : 0.88
4) What is the percentage of correctly and wrongly classified class 0?
   Precision of class 0 : 0.94
5) What is the percentage of correctly and wrongly classified class 1?
   Precision of class 1 : 0.86
6) Measure the balance between precision and recall for class 0?
   F1 score of class 0 : 0.93
7) Measure the balance between precision and recall for class 1?
   F1 score of class 1 : 0.87
8) What is the macro average of precision?
   macro average of precision : 0.90
9) What is the macro average of recall?
   macro average of recall : 0.90
10) What is the macro average of f1 score?
    macro average of f1 score : 0.90
11) What is the weighted average of precision?
    weighted average of precision : 0.91
12) What is the weighted average of recall?
    weighted average of recall : 0.91
13) What is the weighted average of f1 score?
    weighted average of f1 score : 0.91

Algorithm : RandomForestClassifier

Accuracy (overall performance of the model) = 0.91

overall metrics performance (precision,recall,f1-score) = good

Result : **Good Model**

Decision Tree Classifier

Code link: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/2.DecisionTree_classification.ipynb

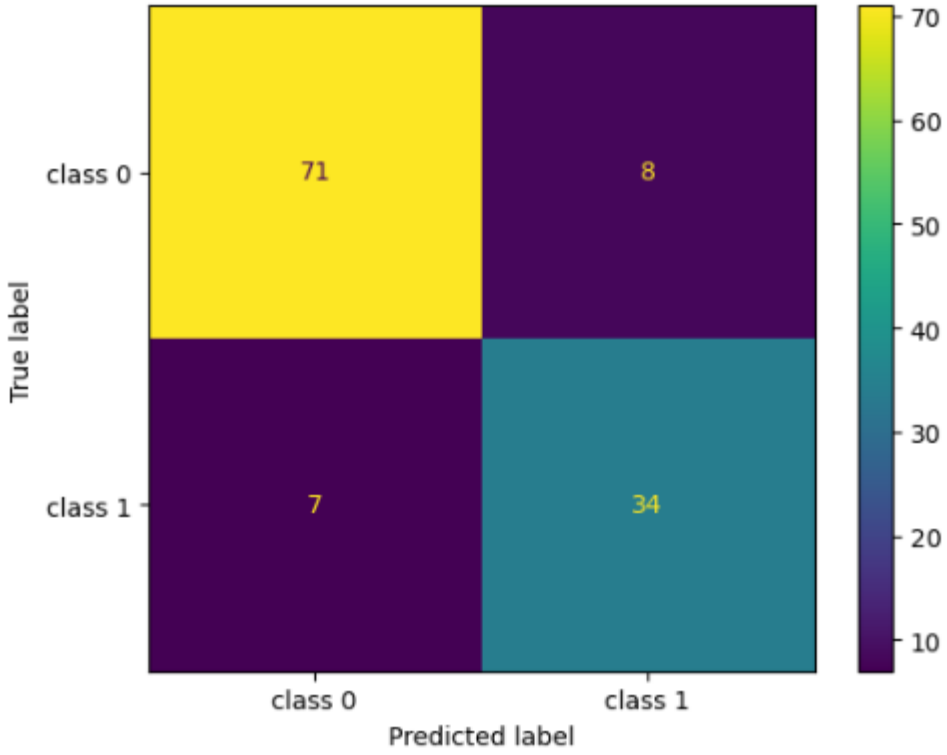

	precision	recall	f1-score	support
0	0.91	0.90	0.90	79
1	0.81	0.83	0.82	41
accuracy			0.88	120
macro avg	0.86	0.86	0.86	120
weighted avg	0.88	0.88	0.88	120

Classification report for random forest classifier

1) What is the percentage of correct classification of both the classes?

   Accuracy : 0.88

2) How many actual positives did I find for class 0?

   Recall of class 0 : 0.90

3) How many actual positives did I find class 1?

   Recall of class 1 : 0.85

4) How correct my positive predictions for class 0?

   Precision of class 0 : 0.92

5) How correct my positive predictions for class 1?
   Precision of class 1 : 0.81
6) What is the overall performance of class 0?
   F1 score of class 0 : 0.93
7) What is the overall performance of class 1?
   F1 score of class 1 : 0.87
8) What is the macro precision?
   macro average of precision : 0.87
9) What is the macro recall?
   macro average of recall : 0.88
10) What is the macro f1 measure?
    macro average of f1 score : 0.87
11) What is the weighted precision?
    weighted average of precision : 0.89
12) What is the weighted recall?
    weighted average of recall : 0.88
13) What is the weighted f1 score?
    weighted average of f1 score : 0.88

Algorithm : DecisionTreeClassifier

Accuracy (overall performance of the model) = 0.88

overall metrics performance (precision,recall,f1-score) = good

Result **: Good Model but not better than Randomforestclassifier**

SVC (support vector classifier)

Code: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/3.SVM_classification.ipynb

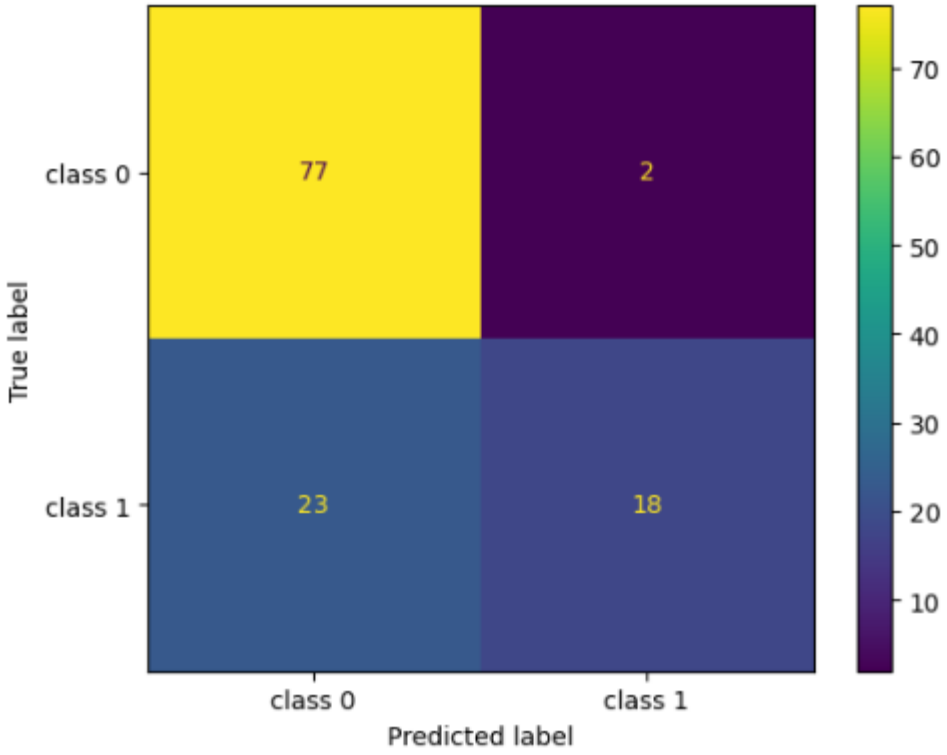

	precision	recall	f1-score	support
0	0.77	0.97	0.86	79
1	0.90	0.44	0.59	41
accuracy			0.79	120
macro avg	0.83	0.71	0.73	120
weighted avg	0.81	0.79	0.77	120

Classification report for random forest classifier

1) Overall how many predictions where correct?

   Accuracy : 0.79

2) Of all actual users for class 0 (Not purchases), How many did the model correctly Identified?

   Recall of class 0 : 0.97

3) Of all actual users for class 1 (purchased), how many did model correctly identified?

   Recall of class 1 : 0.44

4) Of all actual users for class 0 (not purchased), how many were actually correct?

   Precision of class 0 : 0.77

5) Of all actual users for class 1 (purchased), how many were actually correct?

Precision of class 1 : 0.90

6) What is F1 measure of class 0?
   F1 score of class 0 : 0.86
7) What is F1 measure of class 1?
   F1 score of class 1 : 0.59
8) What is the average performance of precision for the model?
   macro average of precision : 0.83
9) What is the average performance of recall for the model?
   macro average of recall : 0.71
10) What is the average performance of f1 score for the model?
    macro average of f1 score : 0.73
11) What is the sum of product of proportion rate of each class in precision?
    weighted average of precision : 0.81
12) What is the sum of product of proportion rate of each class in recall?
    weighted average of recall : 0.79
13) What is the sum of product of proportion rate of each class in f1 score?
    weighted average of f1 score : 0.77

Algorithm : SVC

Accuracy (overall performance of the model) = 0.79

overall metrics performance (precision,recall,f1-score) = poor

Result : **poor model**

LOGISTIC REGRESSION (CLASSIFICATION ALGORITHM)

Code: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/4.Logistic_Regression_classification.ipynb

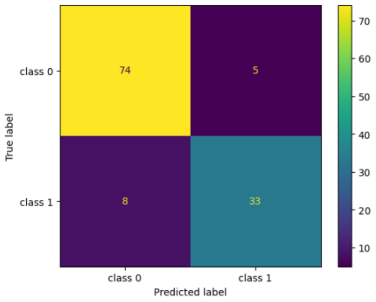

	precision	recall	f1-score	support
0	0.90	0.94	0.92	79
1	0.87	0.80	0.84	41
accuracy			0.89	120
macro avg	0.89	0.87	0.88	120
weighted avg	0.89	0.89	0.89	120

Algorithm : Logistic Regression

Accuracy (overall performance of the model) = 0.89

overall metrics performance (precision,recall,f1-score) = good

Result : **good model but not better than Randomforestclassifier**

PROBLEM STATEMENT:- Social network Ad (User Purchase Prediciton)
Algorithm: Logistic regression

SL.NO	solver	max_iter	C	Confusion matrix	model accuracy	Remark
1	lbfgs	100	1.0	[[74 5] [8 33]]	0.89 (ConvergenceWarning)	Warning
2	lbfgs	120	1.0	[[74 5] [8 33]]	0.89	Good Model
3	lbfgs	200	1.0	[[74 5] [8 33]]	0.89	Good Model
4	liblinear	120	1.0	[[79 0] [41 0]]	0.66	poor
5	newton-cg	120	1.0	[[74 5] [8 33]]	0.89	Good Model
6	newton-cholesky	120	1.0	[[74 5] [8 33]]	0.89	Good Model
7	sag	120	1.0	[[79 0] [41 0]]	0.66	poor
8	saga	120	1.0	[[79 0] [41 0]]	0.66	poor
9	lbfgs	120	0.01	[[74 5] [11 30]]	0.87	Good Model

KNN classifier:

Code: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/4.KNN_classification.ipynb

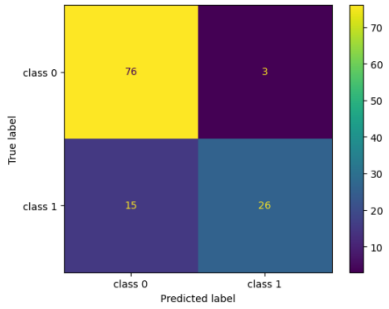

	precision	recall	f1-score	support
0	0.84	0.96	0.89	79
1	0.90	0.63	0.74	41
accuracy			0.85	120
macro avg	0.87	0.80	0.82	120
weighted avg	0.86	0.85	0.84	120

Algorithm : K nearest neighbors classifier

Accuracy (overall performance of the model) = 0.85

overall metrics performance (precision,recall,f1-score) = good

Result**: good model but not better than Randomforestclassifier for the problem statement (social network ad)**

PROBLEM STATEMENT:- Social network Ad (User Purchase Prediciton) Algorithm: K nearest neighbors								
SL.NO	**n_neighbors**	**metric**	**P (power parameter)**	**algorithm**	**weights**	**Confusion matrix**	**model accuracy**	**Remark**
1	5	minkowski	2	auto	uniform	[[69 10] [11 30]]	0.82	poor
2	7	minkowski	2	auto	uniform	[[72 7] [13 28]]	0.83	poor
3	8	minkowski	2	auto	uniform	[[76 3] [16 25]]	0.84	poor
4	11	minkowski	2	auto	uniform	[[74 5] [13 28]]	0.85	Good Model
5	21	minkowski	2	auto	uniform	[[76 3] [15 26]]	0.85	Good Model
6	21	minkowski	2	auto	distance	[[58 21] [10 31]]	0.74	poor
7	21	minkowski	2	ball_tree	uniform	[[76 3] [15 26]]	0.85	Good Model
8	21	minkowski	2	kd_tree	uniform	[[76 3] [15 26]]	0.85	Good Model
9	21	minkowski	2	brute	uniform	[[76 3] [15 26]]	0.85	Good Model

Naïve Bayes

Code: https://github.com/krthiksha/Machine-Learning-Classification_module/blob/main/4.NB_classification.ipynb

PROBLEM STATEMENT:- Social network Ad (User Purchase Prediciton) Algorithm: Naïve Bayes				
SL.NO	**NB types**	**Confusion matrix**	**model accuracy**	**Remark**
1	**GaussianNB**	[[74 5] [8 33]]	0.89	Good Model
2	MultinomialNB	[[68 11] [28 13]]	0.68	poor
3	ComplementNB	[[42 37] [20 21]]	0.53	poor
4	BernoulliNB	[[79 0] [41 0]]	0.66	poor
5	**CategoricalNB**	[[76 3] [9 32]]	0.90	Good Model

Evaluation Metrics using Confusion Matrix

Accuracy

$$\text{Accuracy} = \frac{\text{Correct predictions}}{\text{All predictions}} \qquad \text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Accuracy Scenarios:

- Overall performance of the model?
- What is the Percentage of correct classification of both the classes?
- Overall how many predictions were correct?

Calculation of Random Forest classifier:

Accuracy = (73+36) / (73+36+5+6)

=109/120

=0.90833 ~ 0.91

Recall

$$\text{Recall (or Sensitivity)} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}}$$

Recall Scenarios:

- Percentage of correctly classified of a specific class?
- How many actual positives did I find?
- Of all actual users for a specific class, how many did the model correctly identified?

Calculation of Random Forest classifier:

Recall (class 0) = 73 / (73+6)

= 73/79 = 0.92

Precision

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Precision Scenarios:

- Percentage of correctly and wrongly classified of a specific class?
- How correct my positive predictions for a specific class?
- Of all users for a specific class, how many were actually correct?

Calculation of Random Forest classifier:

Precision (class 0) = 73 / (73+5) = 73/78

= 0.9358 ~ 0.94

F1 score / F1 measure

$$F1\ score = \frac{2}{\frac{1}{Precision} + \frac{1}{Recall}} = 2 \cdot \frac{Precision * Recall}{Precision + Recall}$$

$$\Rightarrow F1\ score = 2 \cdot \frac{Precision * Recall}{Precision + Recall}$$

Scenarios:

- Balance between precision and recall for classes?
- Overall performance of a specific class?
- What is F1 measure of a specific class?

Calculation of Random Forest classifier:

F1 score (class 0) = 2 ((0.94 * 0.92) / (0.94 + 0.92))

= 2(0.8648/1.86)

= 2(0.464)

=0.9298 ~ 0.93

Macro average

Macro precision = (precision1 + precision2) / 2

Macro recall = (recall1 + recall2) / 2

Macro f1score = (f1score 1 + f1score 2) / 2

Calculation of Random Forest classifier:

Macro precision = (0.94 + 0.86) / 2 = 1.8/2 =0.90

Scenarios:

Metrics – (precision, recall, f1 score)

- Macro average of metrics
- Macro metrics
- Average performance for the metrics for the model

Weighted average

Weighted average = Sum (metrics * proportion rate)

Scenarios:

Metrics – (precision, recall, f1 score)

- Weighted average of metrics
- Weighted metrics
- Sum of product of proportional rate of each classes in metrics

Calculation of Random Forest classifier:

Weighted precision = (0.94 * (79/120)) + (0.86 * (41/120))

= (0.94 * (0.658)) + (0.86 * (0.34))

= 0.618 + 0.29

= 0.9085 ~ 0.91